

#### -----WORD MEAN PROGRAM AND EXECUTION-----

**1)Created an input file for the wordmean program in path /hadoopdata/hdfs/in with file name wordmean.txt**

```
[acadgild@localhost ~]$ hadoop fs -mkdir /hadoopdata/hdfs/in
```

```
[acadgild@localhost ~]$ hadoop fs -put wordmean.txt /hadoopdata/hdfs/in
```

```
[acadgild@localhost ~]$ hadoop fs -cat /hadoopdata/hdfs/in/wordmean.txt
```

good morning everyone welcome to the acadgild session specially to the toughest part which is yarn.

**2)Ran the word mean program using the mentioned hadoop-mapreduce-examples-2.7.2.jar in path /home/acadgild/hadoop-2.7.2/share/hadoop/mapreduce using following command and got output has mean 5.25**

```
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.7.2.jar wordmean /hadoopdata/hdfs/in/wordmean.txt /hadoopdata/hdfs/out/filewm
```

Java HotSpot(TM) Client VM warning: You have loaded library /home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.

It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/05/13 21:06:20 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/05/13 21:06:22 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/05/13 21:06:25 INFO input.FileInputFormat: Total input paths to process : 1

18/05/13 21:06:25 INFO mapreduce.JobSubmitter: number of splits:1

18/05/13 21:06:26 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_1526194450906\_0002

18/05/13 21:06:27 INFO impl.YarnClientImpl: Submitted application application\_1526194450906\_0002

18/05/13 21:06:28 INFO mapreduce.Job: The url to track the job:  
[http://localhost:8088/proxy/application\\_1526194450906\\_0002/](http://localhost:8088/proxy/application_1526194450906_0002/)

18/05/13 21:06:28 INFO mapreduce.Job: Running job: job\_1526194450906\_0002

18/05/13 21:06:55 INFO mapreduce.Job: Job job\_1526194450906\_0002 running in uber mode : false

18/05/13 21:06:55 INFO mapreduce.Job: map 0% reduce 0%

18/05/13 21:07:09 INFO mapreduce.Job: map 100% reduce 0%

18/05/13 21:07:26 INFO mapreduce.Job: map 100% reduce 100%

18/05/13 21:07:27 INFO mapreduce.Job: Job job\_1526194450906\_0002 completed successfully

18/05/13 21:07:28 INFO mapreduce.Job: Counters: 49

#### File System Counters

FILE: Number of bytes read=39

FILE: Number of bytes written=235393

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=219

HDFS: Number of bytes written=19

HDFS: Number of read operations=6

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

#### Job Counters

Launched map tasks=1

Launched reduce tasks=1

Data-local map tasks=1

Total time spent by all maps in occupied slots (ms)=11136

Total time spent by all reduces in occupied slots (ms)=13258

Total time spent by all map tasks (ms)=11136

Total time spent by all reduce tasks (ms)=13258

Total vcore-milliseconds taken by all map tasks=11136

Total vcore-milliseconds taken by all reduce tasks=13258

Total megabyte-milliseconds taken by all map tasks=11403264

Total megabyte-milliseconds taken by all reduce tasks=13576192

#### Map-Reduce Framework

Map input records=1

Map output records=32

Map output bytes=464

Map output materialized bytes=39

Input split bytes=118

Combine input records=32

Combine output records=2

Reduce input groups=2

Reduce shuffle bytes=39

Reduce input records=2

Reduce output records=2

Spilled Records=4

Shuffled Maps =1

Failed Shuffles=0

Merged Map outputs=1

GC time elapsed (ms)=499

CPU time spent (ms)=2550

Physical memory (bytes) snapshot=229715968

Virtual memory (bytes) snapshot=635781120

Total committed heap usage (bytes)=137498624

#### Shuffle Errors

BAD\_ID=0

CONNECTION=0

IO\_ERROR=0

WRONG\_LENGTH=0

WRONG\_MAP=0

WRONG\_REDUCE=0

#### File Input Format Counters

Bytes Read=101

#### File Output Format Counters

Bytes Written=19

**The mean is: 5.25**

#### OUTPUT:

```
[acadgild@localhost ~]$ hadoop fs -ls /hadoopdata/hdfs/out/filewm
```

Found 2 items

```
-rw-r--r--  1 acadgild supergroup      0 2018-05-13 21:07 /hadoopdata/hdfs/out/filewm/_SUCCESS
```

```
-rw-r--r--  1 acadgild supergroup    19 2018-05-13 21:07 /hadoopdata/hdfs/out/filewm/part-r-00000
```

```
[acadgild@localhost ~]$ hadoop fs -cat /hadoopdata/hdfs/out/filewm/part-r-00000
```

**count 16**

**length 84**

