----WORDSTANDARDDEVIATION PROGRAM EXECUTION------

1)Created an input file for the wordmean program in path /hadoopdata/hdfs/in with file name wordmean.txt

[acadgild@localhost ~]$ hadoop fs -mkdir /hadoopdata/hdfs/in

[acadgild@localhost ~]$ hadoop fs -put wordmean.txt /hadoopdata/hdfs/in

[acadgild@localhost ~]$ hadoop fs -cat /hadoopdata/hdfs/in/wordmean.txt

good morning everyone welcome to the acadgild session specially to the toughest part whcih is yarn.

2)Ran the word mean program using the mentioned hadoop-mapreduce-examples-2.7.2.jar in path /home/acadgild/hadoop-2.7.2/share/hadoop/mapreduce using following command.

[acadgild@localhost mapreduce]$ **hadoop jar hadoop-mapreduce-examples-2.7.2.jar wordstandarddeviation /hadoopdata/hdfs/in/wordmean.txt /hadoopdata/hdfs/out/fileWordSD**

Java HotSpot(TM) Client VM warning: You have loaded library /home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.

It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.

18/05/13 23:34:48 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/05/13 23:34:50 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

18/05/13 23:34:52 INFO input.FileInputFormat: Total input paths to process : 1

18/05/13 23:34:53 INFO mapreduce.JobSubmitter: number of splits:1

18/05/13 23:34:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1526194450906_0006

18/05/13 23:34:54 INFO impl.YarnClientImpl: Submitted application application_1526194450906_0006

18/05/13 23:34:54 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1526194450906_0006/

18/05/13 23:34:54 INFO mapreduce.Job: Running job: job_1526194450906_0006

18/05/13 23:35:19 INFO mapreduce.Job: Job job_1526194450906_0006 running in uber mode : false

18/05/13 23:35:19 INFO mapreduce.Job: map 0% reduce 0%

18/05/13 23:35:34 INFO mapreduce.Job: map 100% reduce 0%

18/05/13 23:35:47 INFO mapreduce.Job: map 100% reduce 100%

18/05/13 23:35:48 INFO mapreduce.Job: Job job_1526194450906_0006 completed successfully

18/05/13 23:35:49 INFO mapreduce.Job: Counters: 49

    File System Counters

        FILE: Number of bytes read=56

        FILE: Number of bytes written=235603

        FILE: Number of read operations=0

        FILE: Number of large read operations=0

        FILE: Number of write operations=0

        HDFS: Number of bytes read=219

        HDFS: Number of bytes written=30

        HDFS: Number of read operations=6

        HDFS: Number of large read operations=0

        HDFS: Number of write operations=2

    Job Counters

        Launched map tasks=1

        Launched reduce tasks=1

        Data-local map tasks=1

        Total time spent by all maps in occupied slots (ms)=11141

        Total time spent by all reduces in occupied slots (ms)=9995

        Total time spent by all map tasks (ms)=11141

Total time spent by all reduce tasks (ms)=9995

Total vcore-milliseconds taken by all map tasks=11141

Total vcore-milliseconds taken by all reduce tasks=9995

Total megabyte-milliseconds taken by all map tasks=11408384

Total megabyte-milliseconds taken by all reduce tasks=10234880

Map-Reduce Framework

Map input records=1

Map output records=48

Map output bytes=704

Map output materialized bytes=56

Input split bytes=118

Combine input records=48

Combine output records=3

Reduce input groups=3

Reduce shuffle bytes=56

Reduce input records=3

Reduce output records=3

Spilled Records=6

Shuffled Maps =1

Failed Shuffles=0

Merged Map outputs=1

GC time elapsed (ms)=514

CPU time spent (ms)=2510

Physical memory (bytes) snapshot=229175296

Virtual memory (bytes) snapshot=635781120

Total committed heap usage (bytes)=137498624

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Input Format Counters

Bytes Read=101

File Output Format Counters

Bytes Written=30

The standard deviation is: 2.384848003542364

OUTPUT:

[acadgild@localhost mapreduce]$ **hadoop fs -ls /hadoopdata/hdfs/out/fileWordSD**

Found 2 items

-rw-r--r--   1 acadgild supergroup   0 2018-05-13 23:35 /hadoopdata/hdfs/out/fileWordSD/_SUCCESS

-rw-r--r--   1 acadgild supergroup   30 2018-05-13 23:35 /hadoopdata/hdfs/out/fileWordSD/part-r-00000

[acadgild@localhost mapreduce**]$ hadoop fs -cat /hadoopdata/hdfs/out/fileWordSD/part-r-00000**

count   16

length   84

square   532